

Visualization of Inferred Versioned Schemas from NoSQL Databases*

Alberto Hernández Chillón

alberto.hernandez1@um.es

Cátedra SAES

University of Murcia

Diego Sevilla Ruiz

dsevilla@ditec.um.es

Faculty of Computer Science

University of Murcia

Jesús García Molina

jmolina@um.es

Faculty of Computer Science

University of Murcia

September 20, 2016

Abstract

While the concept of database schema plays a central role in Relational Database systems, most NoSQL systems are schemaless: these databases are created without having to formally define its schema, that remains implicit in the stored data. This lack of schema definition offers a greater flexibility; more specifically, it eases both the recording of non-uniform data and data evolution. However, this comes at the cost of losing some of the benefits provided by schemas (such as type checking and integrity checking). In [1], a MDE-based reverse engineering approach for inferring the schema of aggregate-oriented NoSQL databases is presented. The obtained schemas can be used to build database utilities that tackle some of the problems encountered using implicit schemas. However, in the NoSQL context, the definition of a unique schema is not enough to accurately describe the stored data: different versions of the same object type coexist in the database, leading to different schema versions. In this work¹ a modeling transformation chain is designed in order to obtain these schema versions from an input database, and a visualization tool for these schemas is defined using Sirius. Sirius provides us with a workbench integrated in Eclipse able to associate graphical representations to metamodel elements, and the automatic generation of an embedded editor in Eclipse, all of this

*Work partially supported by the Cátedra SAES of the University of Murcia (<http://www.catedrasaes.org>), a research lab sponsored by the SAES company (<http://www.electronica-submarina.com/>).

¹A complete description of this M.Sc. project may be found on the GitHub repository project: <https://github.com/Soltari/NoSQLVisualizationTools>.

with a suitable distribution method such as plugins and update sites. The visualization tool consists of several viewpoints according to our requirements. The most important designed viewpoint is the tree viewpoint. This representation allows the user to navigate between domain concepts such as *Schema version*, *Entity*, and *Entity version*, and to analyze these concepts with more viewpoints. The tree viewpoint also provides an inverted index and a dictionary. From this viewpoint it is possible to navigate to other diagrams such as the global schema viewpoint, which shows all the *Entities*, *Entity versions* and *Properties* in an intuitive way. Now for each inferred version schema, a couple of viewpoints are provided to analyze all the objects contained on the schema. Finally, a simple viewpoint is designed to inspect standalone *Entities*.

References

- [1] Diego Sevilla Ruiz, Severino Feliciano Morales, and Jesús García Molina. Inferring Versioned Schemas from NoSQL Databases and its Applications. In *ER 2015*, pages 467–480, September 2015.

Additional Figures




Figure 1: The process is composed of the definition of a *NoSQL_Schema* metamodel and an *Extended_NoSQL_Schema* metamodel to identify *Schema versions*. With a *m2m* transformation and some input models from the first metamodel, models of the second metamodel are generated. Now with Sirius these models may be visualized with several viewpoints.

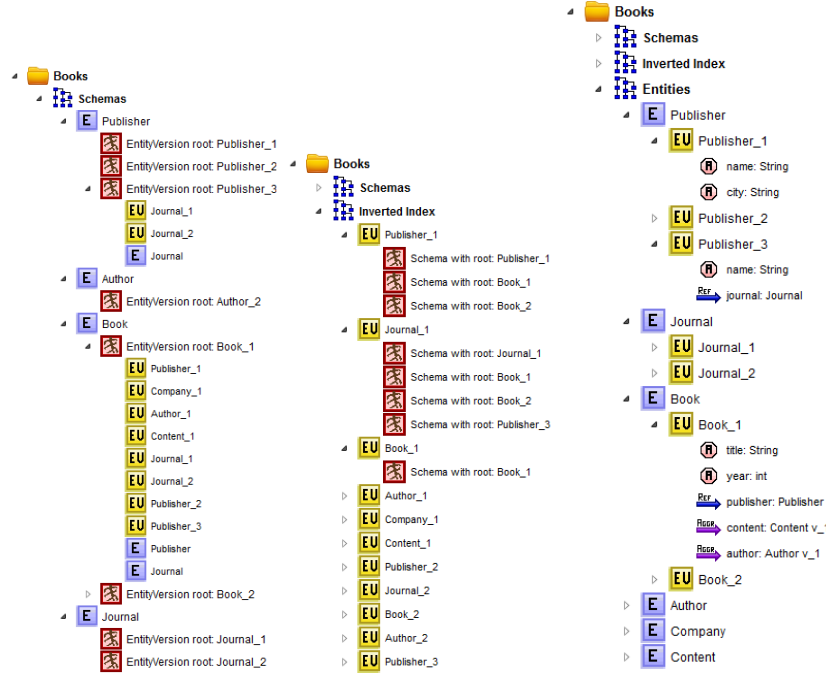


Figure 2: This tree viewpoint is divided into three folders. The first one stores *Schema versions* sorted by *Entity*. The second folder stores *Entity versions*. The last folder contains a dictionary of *Entities*, *Entity Versions* and *Properties*. Each element may be expanded to show its components.

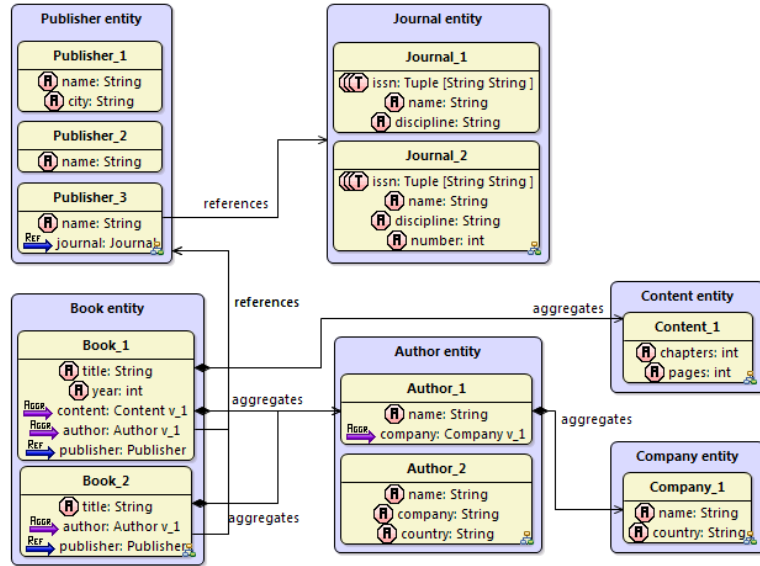


Figure 3: This global schema viewpoint shows a diagram with several box-like entities, nested elements (aggregates), elements properties and some connections. Each one of the elements is represented with a distinctive icon.

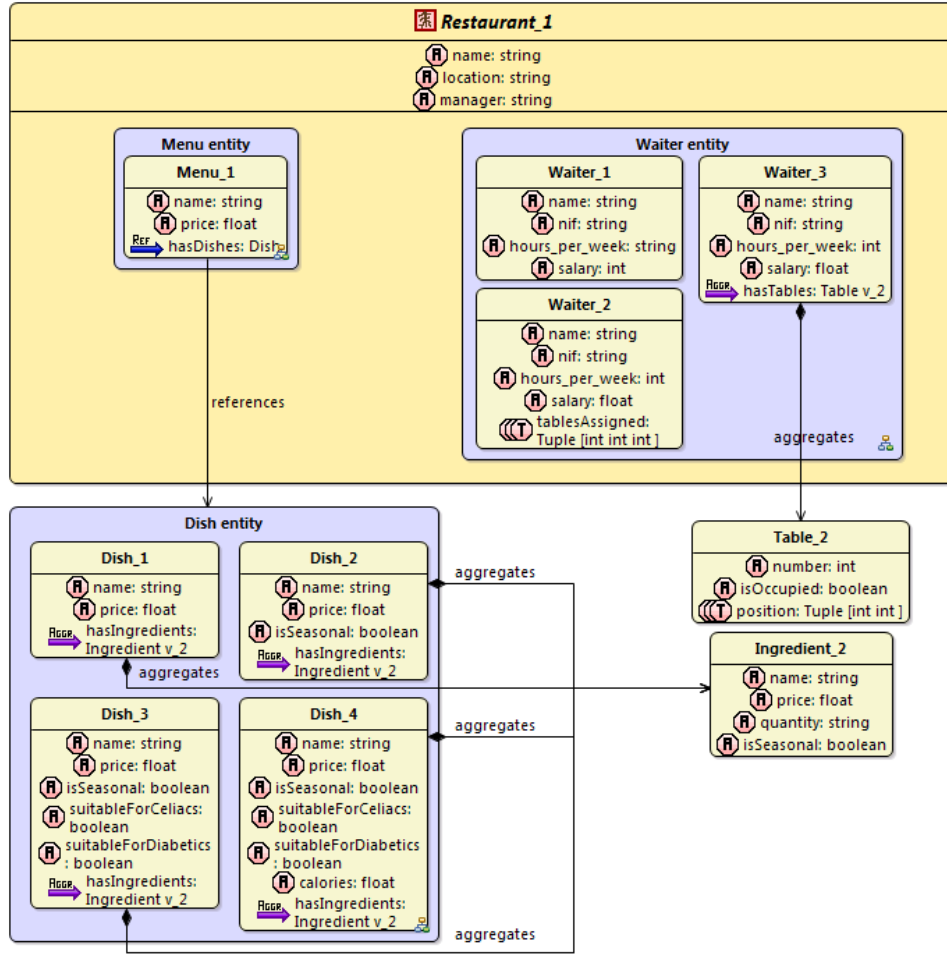


Figure 4: This specific schema viewpoint shows all the elements which compose a *Schema version*. A *Schema version* contains an *Entity Version* root, several *Entities* and *Entity versions*. The same distinctive icons are used for all viewpoints.

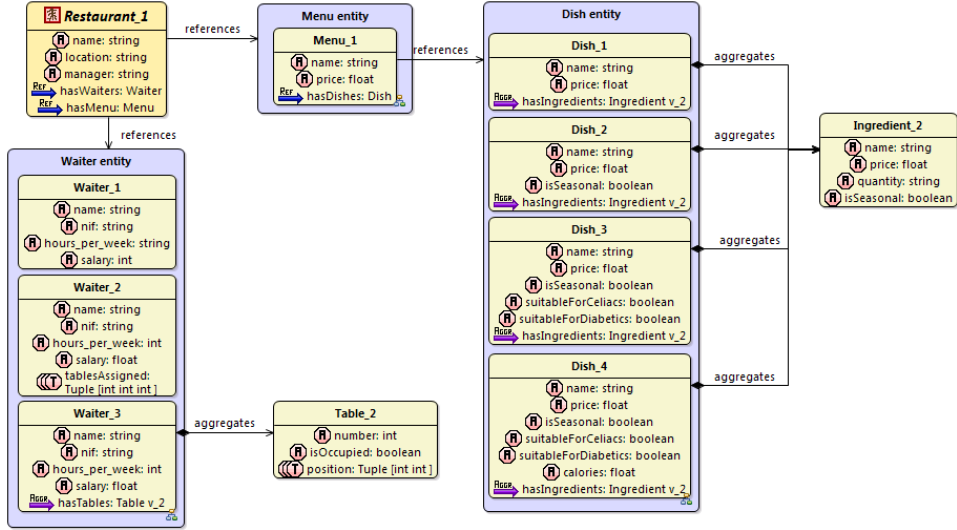


Figure 5: This schema viewpoint is equivalent to the last one but changing the way elements are nested. It organizes the elements with the same layout as the global schema viewpoint.

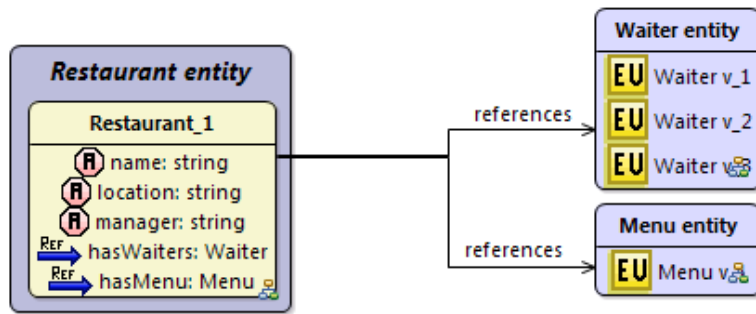


Figure 6: This last viewpoint shows a detailed *Entity* with all its *Entity versions* and all the connections with other elements that these *Entity versions* have.